


The background features a dynamic, abstract design with flowing, wavy lines in shades of blue and white, creating a sense of movement and depth. The lines are more concentrated on the right side, where they form a dense, swirling pattern, and become more sparse and lighter towards the left.

Operating System Review (Chapter 3-5)

3.1 Process Concept

- What are the most important four components of a process?
- ❖ Text section: program code itself
- ❖ Stack: temporary data (function parameters, return addresses, local variables)
- ❖ Data section: global variables
- ❖ Heap: contains memory dynamically allocated during run-time

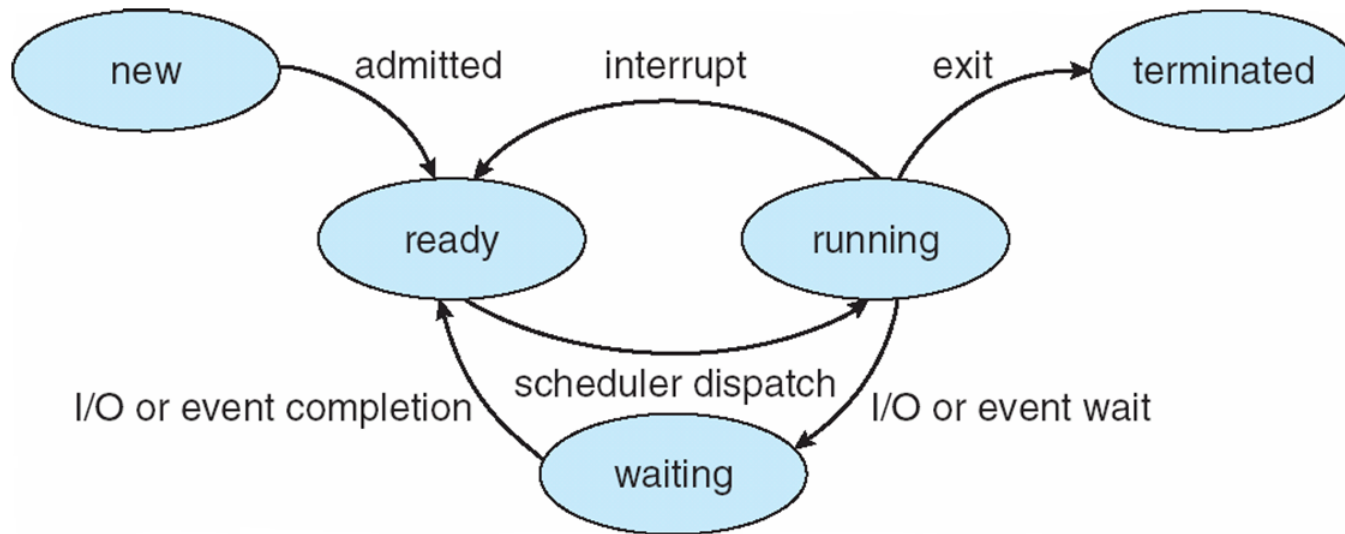
max



0

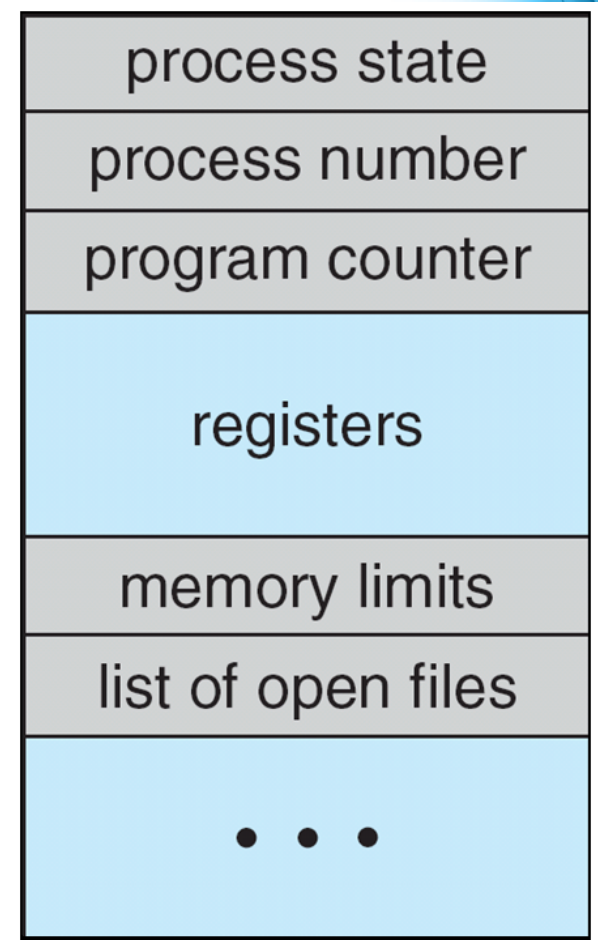
3.1 Process Concept

- Provide at least three possible states a process may be in



3.1 Process Concept

- What is a Process Control Block (PCB) ?
- ❖ Process state – running, waiting, etc
- ❖ Program counter – location of instruction to next execute
- ❖ CPU registers – contents of all process-centric registers
- ❖ CPU scheduling information- priorities, scheduling queue pointers
- ❖ Memory-management information – memory allocated to the process
- ❖ Accounting information – CPU used, clock time elapsed since start, time limits
- ❖ I/O status information – I/O devices allocated to process, list of open files



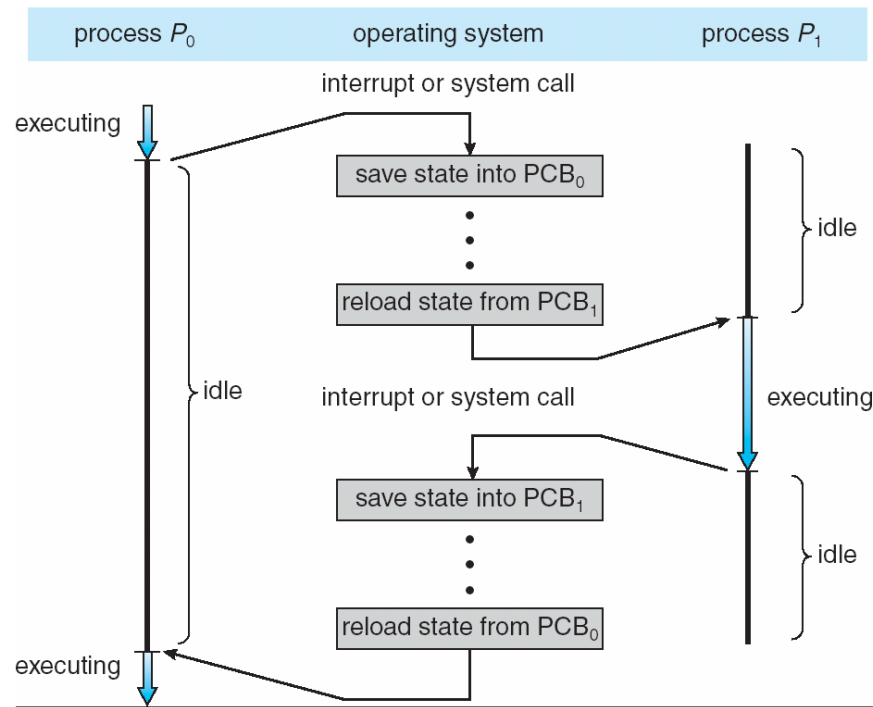
3.1 Process Concept

- List two process types

- ❖ I/O Bound: spends more time doing I/O than computations, many short CPU bursts
- ❖ CPU Bound: spends more time doing computations, few very long CPU bursts

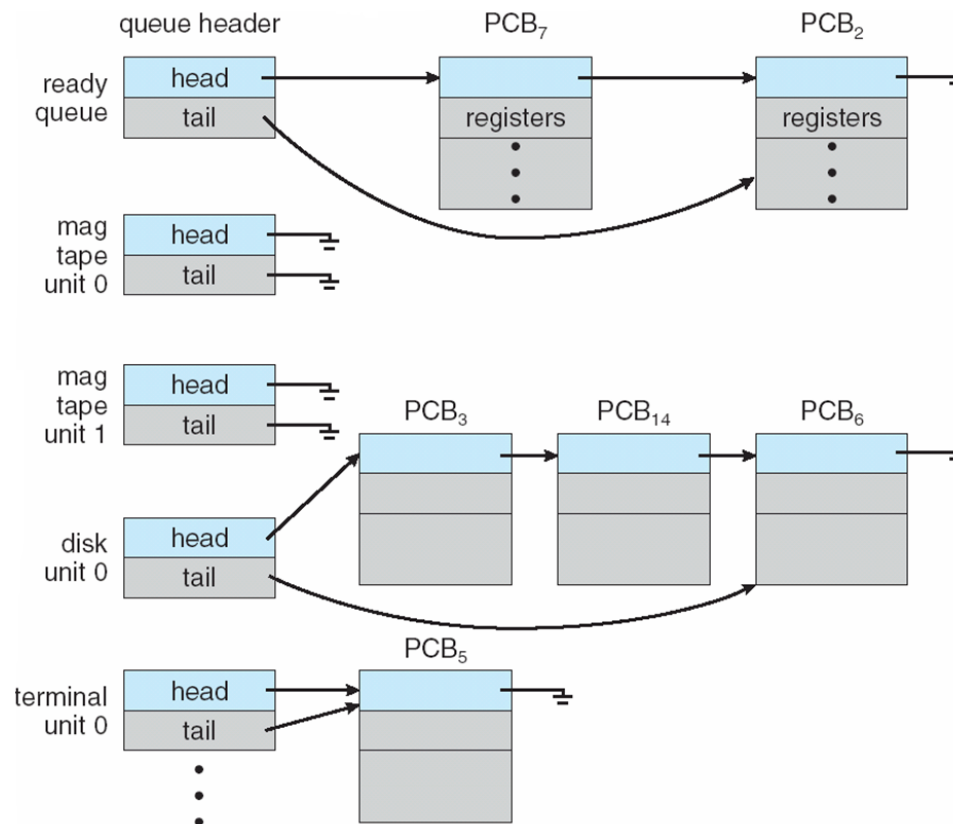
3.2 Process Scheduling

- What is the term that describes saving the state of one process, and restoring the state of another?
 - ❖ When CPU switches to another process, the system must save the state of the old process (to PCB) and load the saved state (from PCB) for the new process via a **context switch**



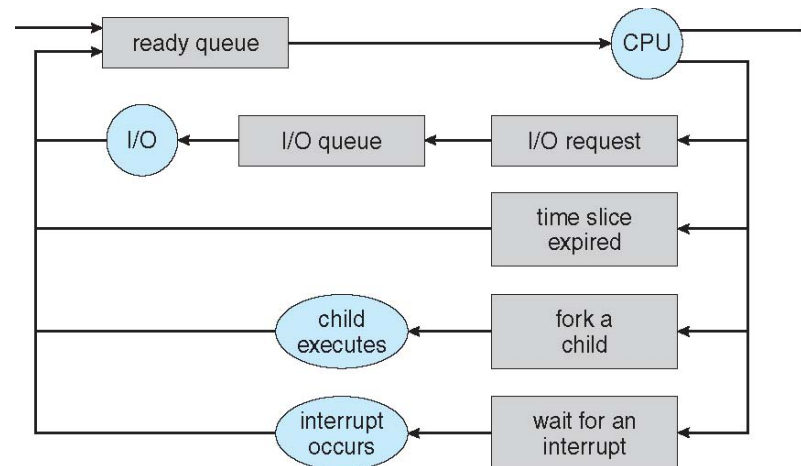
3.2 Process Scheduling

- How many queues for the process scheduler?
 - ❖ **Job queue** – set of all processes in the system
 - ❖ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - ❖ **Device queues** – set of processes waiting for an I/O device



3.2 Process Scheduling

- Show the difference between short-term and long-term scheduling?
- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue



3.3 Process Creation

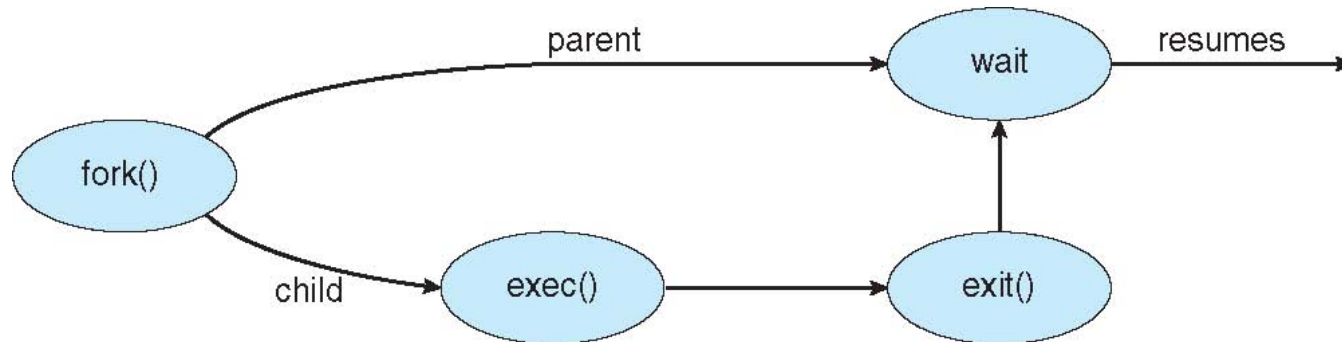
- How to create a new process?
 - ❖ **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
 - ❖ Generally, process identified and managed via a **process identifier (pid)**
 - ❖ Resource sharing options
 - ❖ Parent and children share all resources
 - ❖ Children share subset of parent's resources
 - ❖ Parent and child share no resources

3.3 Process Creation

- What is the difference between fork and exec?

fork() system call creates new process

exec() system call used after a **fork()** to replace the process' memory space with a new program



3.3 Process Creation

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

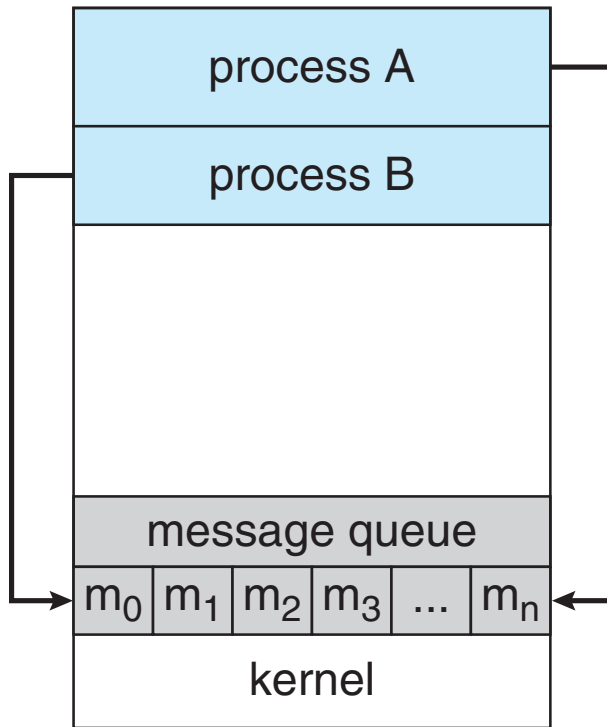
    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }


    return 0;
}
```

3.4 Process Communication

- What are the two models of IPC (**inter process communication**)?
 - **Shared memory**
 - **Message passing**



(a)



(b)

3.4 Process Communication

- What are the two models of message passing?

- ❖ Direct Communication

- **send** (*P, message*) – send a message to process P
 - **receive**(*Q, message*) – receive a message from process Q

- ❖ Indirect Communication

- **send**(*A, message*) – send a message to mailbox A
 - **receive**(*A, message*) – receive a message from mailbox A

3.4 Process Communication

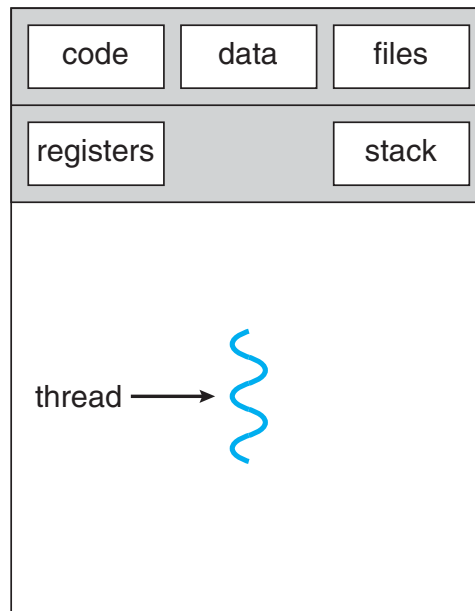
- What is the difference between blocking model and non-blocking model?
 - ❖ **Blocking** is considered **synchronous**
 - ❖ **Blocking send** -- the sender is blocked until the message is received
 - ❖ **Blocking receive** -- the receiver is blocked until a message is available
 - ❖ **Non-blocking** is considered **asynchronous**
 - ❖ **Non-blocking send** -- the sender sends the message and continue
 - ❖ **Non-blocking receive** -- the receiver receives:
 - ❖ A valid message, or
 - ❖ Null message

3.4 Process Communication

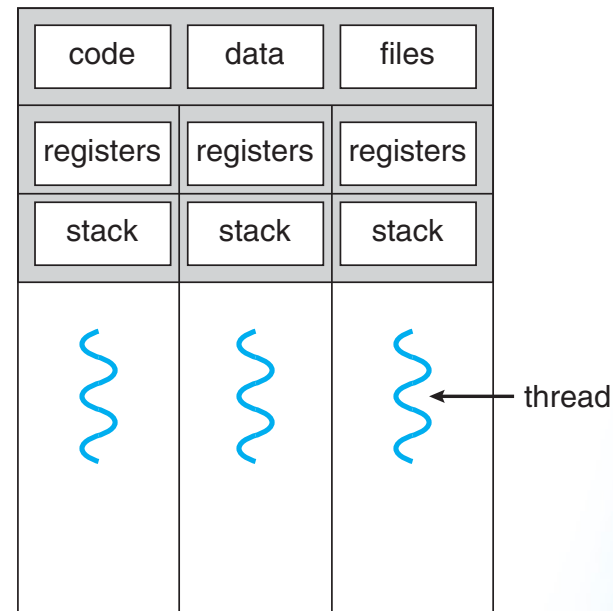
- What are the most popular three methods to communicate between client and server?
 - ❖ Sockets
 - ❖ Remote Procedure Calls
 - ❖ Remote Method Invocation (Java)

4.1 Thread Concept

- What data are shared between threads inside the same process?
 - ❖ Code
 - ❖ Data
 - ❖ Files



single-threaded process



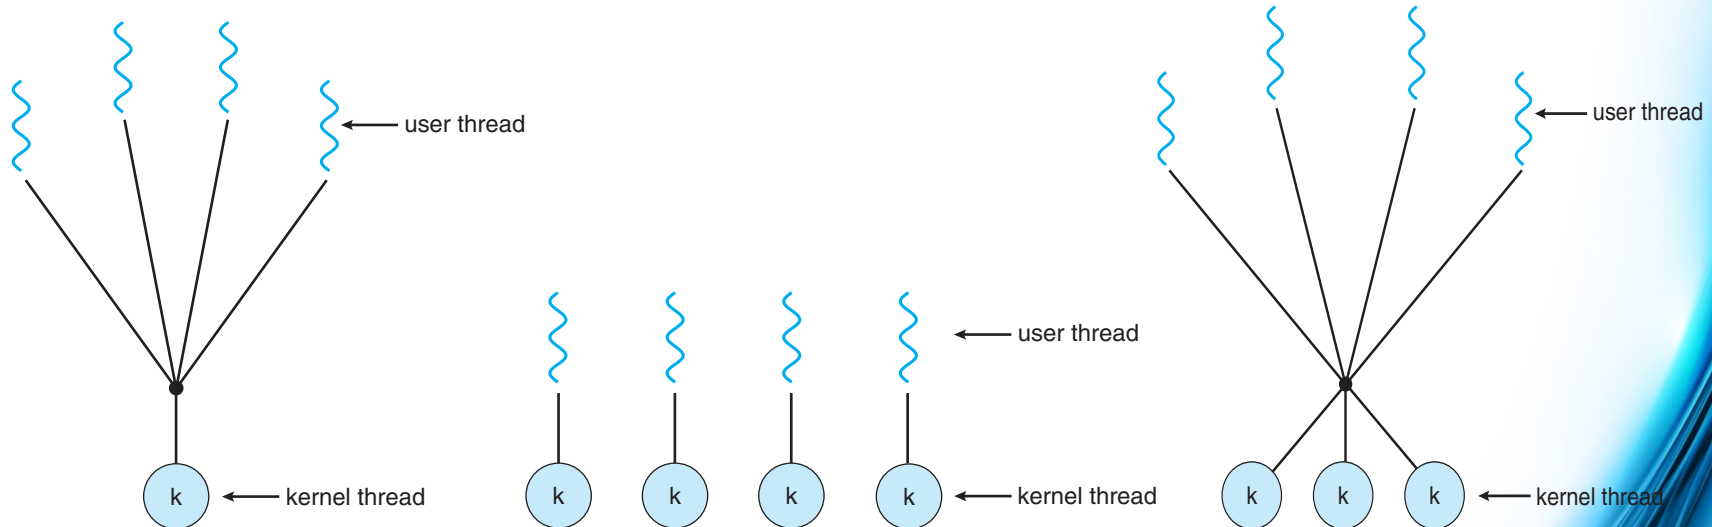
multithreaded process

4.1 Thread Concept

- Provide at least three benefits of multithreaded programming
 - ❖ **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
 - ❖ **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
 - ❖ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
 - ❖ **Scalability** – process can take advantage of multiprocessor architectures

4.1 Thread Concept

- What are the three basic mapping relationship between kernel thread and user thread?
 - ❖ Many-to-One
 - ❖ One-to-One
 - ❖ Many-to-Many




4.1 Thread Concept

- What is the LWP (lightweight process)?

Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**

- Appears to be a virtual processor on which process can schedule user thread to run
- Each LWP attached to kernel thread
- How many LWPs to create?



4.2 Thread Creation

○ Pthread

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```


4.2 Thread Creation

○ Windows

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}
```

4.2 Thread Creation

○ Java

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```

4.2 Thread Creation

- Why do we need thread pool?

- ❖ Create a number of threads in a pool where they await work

- ❖ Advantages:

- ❖ Usually slightly faster to service a request with an existing thread than create a new thread

- ❖ Allows the number of threads in the application(s) to be bound to the size of the pool

- ❖ Separating task to be performed from mechanics of creating task allows different strategies for running task

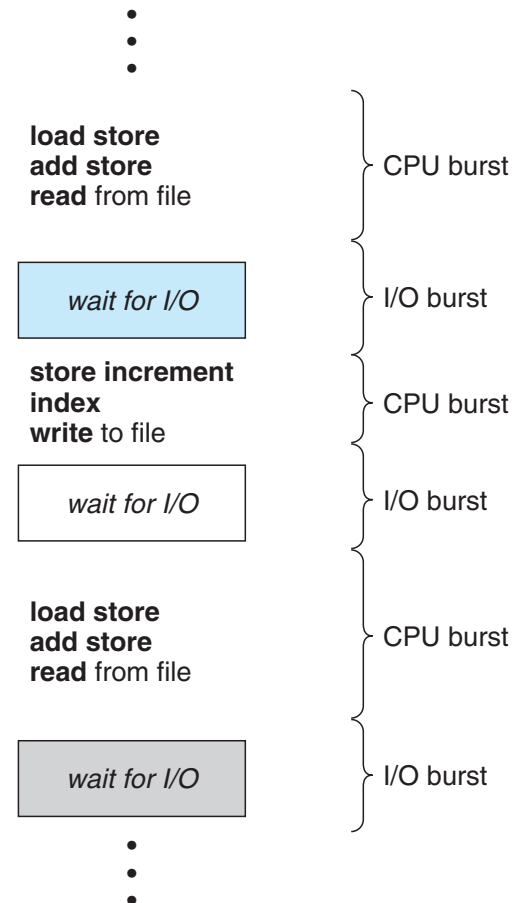
- ❖ i.e.Tasks could be scheduled to run periodically

4.2 Thread Creation

- What are the two models of thread cancellation?
 - Asynchronous cancellation terminates the target thread immediately
 - Deferred cancellation allows the target thread to periodically check if it should be canceled

5.1 CPU Scheduling

- What are the two bursts that CPU schedulers are designed around?
 - CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
 - **CPU burst** followed by **I/O burst**



5.1 CPU Scheduling

- True or False? Under preemptive scheduling, when a process switches from the running to the ready state, it may lose control of the CPU.
 - Switches from running to waiting (nonpreemptive)
 - Switches from running to ready (preemptive)
 - Switches from waiting to ready (preemptive)
 - Terminates (nonpreemptive)

5.2 Scheduling Algorithm

- List at least three different criteria for designing a CPU scheduling algorithm
 - **CPU utilization** – keep the CPU as busy as possible
 - **Throughput** – # of processes that complete their execution per time unit
 - **Turnaround time** – amount of time to execute a particular process
 - **Waiting time** – amount of time a process has been waiting in the ready queue
 - **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

5.2 Scheduling Algorithm

- What scheduling algorithm assigns the CPU to the process that first requested it?
 - First- Come, First-Served (FCFS) Scheduling
- What scheduling algorithm assigns the CPU to the process with the shortest burst?
 - First Job Scheduling (FJS)
- What scheduling algorithm assigns the CPU to a process for only its time slice (or time quantum?)
 - round robin scheduling (RR)

5.2 Scheduling Algorithm

- What scheduling algorithm assigns the CPU to the process with the highest priority?
 - Priority Scheduling
- True or false, Shortest Job First is a specific priority scheduling algorithm?
- True or False? The multilevel feedback queue scheduling algorithm allows processes to migrate between different queues.