

# Lab 0: RV64内核调试

## 1. 实验目的

安装虚拟机及Docker，通过在QEMU模拟器上运行Linux来熟悉如何从源代码开始将内核运行在QEMU模拟器上，学习使用GDB跟QEMU对代码进行联合调试，为后续实验打下基础。

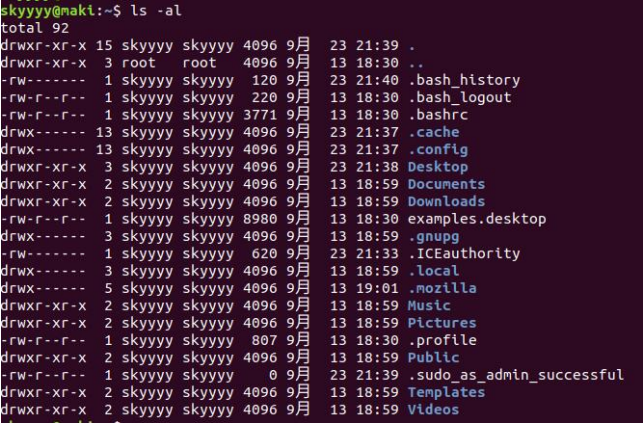
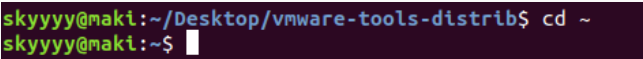
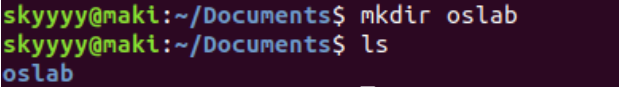
## 2. 实验内容及要求

- 安装虚拟机软件、Ubuntu镜像，自行学习Linux基础命令。
- 安装Docker，下载并导入Docker镜像，创建并运行容器。
- 编译内核并用 gdb + QEMU 调试，在内核初始化过程中设置断点，对内核的启动过程进行跟踪，并尝试使用 gdb的各项命令。

## 3. 操作方法和实验步骤

### 3.1 通过虚拟机安装Linux系统

终端命令

命令	作用	截图
pwd	打印出当前工作目录的名称	
ls	打印出当前工作目录的所有内容	
ls -al	使用长格式打印包括隐藏文件的所有文件及目录信息	
cd ~	切换至当前用户目录	
mkdir oslab	在当前工作目录下创建名为oslab的目录	

vi test.c	用vi编辑器创建并进入名为test.c的文件(vi编辑器命令在此不列)	
gedit test.c	用gedit编辑器打开test.c文件	
rm test.c	删除test.c文件	
sudo apt install curl	连接服务器安装curl	
touch a.txt	在当前工作目录创建a.txt文件	
cat a.txt   tail -n 10	查看a.txt尾部10行的内容	

## 3.2 安装Docker环境并创建容器

### 1. 安装docker

### 使用官方安装脚本自动安装docker

```
$ curl -fsSL https://get.docker.com | bash -s docker --mirror Aliyun
```

### 将用户加入docker组, 免 sudo

```
$ sudo usermod -aG docker $USER    ### 注销后重新登陆生效
```

```
=====
To run Docker as a non-privileged user, consider setting up the
Docker daemon in rootless mode for your user:

    dockerd-rootless-setuptool.sh install

Visit https://docs.docker.com/go/rootless/ to learn about rootless mode.

To run the Docker daemon as a fully privileged service, but granting non-root
users access, refer to https://docs.docker.com/go/daemon-access/

WARNING: Access to the remote API on a privileged Docker daemon is equivalent
to root access on the host. Refer to the 'Docker daemon attack surface'
documentation for details: https://docs.docker.com/go/attack-surface/
=====

skyyyy@naki:~$
```

## 2. 下载并导入docker镜像

### 首先进入oslab.tar所在的文件夹, 然后使用该命令导入docker镜像

```
$ cat oslab.tar | docker import - oslab:2020
```

### 执行命令后若出现以下错误提示

### ERROR: Got permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock

### 可以使用下面命令为该文件添加权限来解决

```
### $ sudo chmod a+rw /var/run/docker.sock
```

### 查看docker镜像

```
$ docker image ls
```

```
skyyyy@naki:~/Documents/oslab$ sudo chmod a+rw /var/run/docker.sock
[sudo] password for skyyyy:
skyyyy@naki:~/Documents/oslab$ cat oslab.tar | docker import - oslab:2020
sha256:6753703b0ae895107da6dd5d702a3aa3a3495c729ad0668731cbf402a3ccd2d6
skyyyy@naki:~/Documents/oslab$ docker image ls
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
oslab         2020      6753703b0ae8   About a minute ago   2.89GB
```

## 3. 从镜像中创建一个容器并进入该容器

### 从镜像创建一个容器

```
$ docker run -it oslab:2020 /bin/bash
```

### -i:交互式操作 -t:终端

```
root@6a50e217856f:/#
```

### 提示符变为 '#' 表明成功进入容器 后面的字符串根据容器而生成, 为容器id

```
root@6a50e217856f:/# exit (或者CTRL+D)
```

### 从容器中退出 此时运行docker ps, 运行容器的列表为空

```
skyyyy@maki:~/Documents/oslab$ docker run -it oslab:2020 /bin/bash
root@6a50e217856f:/#
root@6a50e217856f:/# exit
exit
```

### 查看当前运行的容器

```
$ docker ps
```

### 查看所有存在的容器

```
$ docker ps -a
```

```
skyyyy@maki:~/Documents/oslab$ docker ps
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
skyyyy@maki:~/Documents/oslab$ docker ps -a
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
6a50e217856f   oslab:2020  "/bin/bash"   3 minutes ago   Exited (0) 3 minutes ago           objective_dirac
```

### 启动处于停止状态的容器

```
$ docker start 6a50      ### 6a50 为容器id的前四位，id开头的几位便可标识一个容器
```

```
skyyyy@maki:~/Documents/oslab$ docker start 6a50
6a50
skyyyy@maki:~/Documents/oslab$ docker ps
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
6a50e217856f   oslab:2020  "/bin/bash"   5 minutes ago   Up 10 seconds           objective_dirac
```

### 进入已经运行的容器 oslab的密码为2020

```
$ docker exec -it -u oslab -w /home/oslab 36 /bin/bash
```

```
$ docker exec -it 6a50 /bin/bash
```

```
skyyyy@maki:~/Documents/oslab$ docker exec -it -u oslab -w /home/oslab 6a /bin/bash
oslab@6a50e217856f:~$
```

```
skyyyy@maki:~$ docker exec -it 6a50 /bin/bash
root@6a50e217856f:/#
root@6a50e217856f:/# pwd
/
root@6a50e217856f:/# cd /home/oslab
root@6a50e217856f:/home/oslab# ls
lab0
root@6a50e217856f:/home/oslab# cd lab0
root@6a50e217856f:/home/oslab/lab0# export TOP=`pwd`
root@6a50e217856f:/home/oslab/lab0# mkdir -p build/linux
root@6a50e217856f:/home/oslab/lab0# make -C linux O=$TOP/build/linux \
> CROSS_COMPILE=riscv64-unknown-linux-gnu- \
> ARCH=riscv CONFIG_DEBUG_INFO=y \
> defconfig all -j$(nproc)
```

- `docker run -it oslab:2020 /bin/bash`

`docker run` 指创建一个新的容器并运行一个命令，`-it` 指以交互模式运行容器并为容器分配一个伪输入终端，`oslab:2020` 为使用的镜像，`/bin/bash` 指在容器内执行 `/bin/bash` 命令

- `docker exec -it -u oslab -w /home/oslab 36 /bin/bash`

`docker exec` 指在运行的容器中执行命令, `-it` 指以交互模式运行容器, `-u` 为登陆的用户名, `oslab -w /home/oslab` 指定工作目录为 `/home/oslab`, `36` 为容器id前2位, `/bin/bash` 指在容器内执行 `/bin/bash` 命令

### 3.3 编译linux内核

```
# pwd
/home/oslab
# cd lab0
# export TOP=`pwd`
# mkdir -p build/linux
# make -C linux O=$TOP/build/linux \
    CROSS_COMPILE=riscv64-unknown-linux-gnu- \
    ARCH=riscv CONFIG_DEBUG_INFO=y \
    defconfig all -j$(nproc)
```

```
skyyyy@maki:~$ docker exec -it 6a50 /bin/bash
root@6a50e217856f:/#
root@6a50e217856f:/# pwd
/
root@6a50e217856f:/# cd /home/oslab
root@6a50e217856f:/home/oslab# ls
lab0
root@6a50e217856f:/home/oslab# cd lab0
root@6a50e217856f:/home/oslab/lab0# export TOP=`pwd`
root@6a50e217856f:/home/oslab/lab0# mkdir -p build/linux
root@6a50e217856f:/home/oslab/lab0# make -C linux O=$TOP/build/linux \
> CROSS_COMPILE=riscv64-unknown-linux-gnu- \
> ARCH=riscv CONFIG_DEBUG_INFO=y \
> defconfig all -j$(nproc)
```

```

CC      drivers/gpu/drm/drm_syncobj.o
CC      drivers/gpu/drm/drm_lease.o
CC      drivers/gpu/drm/drm_writeback.o
CC      drivers/gpu/drm/drm_client.o
AR      drivers/of/built-in.a
CC      drivers/gpu/drm/drm_client_modeset.o
CC      drivers/gpu/drm/drm_atomic_uapi.o
CC      drivers/gpu/drm/drm_hdcp.o
CC      drivers/gpu/drm/drm_managed.o
CC      drivers/gpu/drm/drm_gem_shmem_helper.o
CC      drivers/gpu/drm/drm_panel.o
CC      drivers/mmc/core/queue.o
CC      drivers/gpu/drm/drm_of.o
CC      drivers/gpu/drm/drm_pci.o
CC      drivers/gpu/drm/drm_debugfs.o
CC      drivers/gpu/drm/drm_debugfs_crc.o
CC      drivers/gpu/drm/drm_panel_orientation_quirks.o
AR      fs/built-in.a
AR      drivers/mmc/core/built-in.a
AR      drivers/mmc/built-in.a
AR      drivers/gpu/drm/built-in.a
AR      drivers/gpu/built-in.a
AR      drivers/built-in.a
GEN      .version
CHK      include/generated/compile.h
LD      vmlinux.o
MODPOST  vmlinux.symvers
MODINFO  modules.builtin.modinfo
GEN      modules.builtin
LD      .tmp_vmlinux.kallsyms1
KSYM     .tmp_vmlinux.kallsyms1.o
LD      .tmp_vmlinux.kallsyms2
KSYM     .tmp_vmlinux.kallsyms2.o
LD      vmlinux
SYSMAP   System.map
MODPOST  Module.symvers
OBJCOPY  arch/riscv/boot/Image
CC [M]   fs/nfs/flexfilelayout/nfs_layout_flexfiles.mod.o
GZIP     arch/riscv/boot/Image.gz
LD [M]   fs/nfs/flexfilelayout/nfs_layout_flexfiles.ko
Kernel: arch/riscv/boot/Image.gz is ready
make[1]: Leaving directory '/home/oslab/lab0/build/linux'
make: Leaving directory '/home/oslab/lab0/linux'
root@6a50e217856f:/home/oslab/lab0#

```

### 3.4 使用QEMU运行内核

```

# qemu-system-riscv64 -nographic -machine virt -kernel
build/linux/arch/riscv/boot/Image \
-device virtio-blk-device,drive=hd0 -append "root=/dev/vda ro console=ttyS0" \
-bios default -drive file=rootfs.ext4,format=raw,id=hd0 \
-netdev user,id=net0 -device virtio-net-device,netdev=net0

```



```

root@6a50e217856f:/home/oslab/lab0# qemu-system-riscv64 -nographic -machine virt -kernel build/linux/arch/riscv/boot/Image \
> -device virtio-blk-device,drive=hd0 -append "root=/dev/vda ro console=ttyS0" \
> -bios default -drive file=rootfs.ext4,format=raw,id=hd0 \
> -netdev user,id=net0 -device virtio-net-device,netdev=net0

OpenSBI v0.6

      _ _ _ _ _
     | | | | |
    _||_|_|_|_
   | | | | | |
  _||_|_|_|_
 | | | | | |
_|_|_|_|_|_|

Platform Name       : QEMU Virt Machine
Platform HART Features : RV64ACDFIMSU
Platform Max HARTs   : 8
Current Hart        : 0
Firmware Base       : 0x80000000
Firmware Size       : 120 KB
Runtime SBI Version  : 0.2

MIDELEG : 0x0000000000000222
MEDELEG : 0x000000000000b109
PMP0    : 0x0000000080000000-0x000000008001ffff (A)
PMP1    : 0x0000000000000000-0xffffffff (A, D, W, X)

```

```

Starting network: udhcpd: started, V1.31.1
udhcpd: sending discover
udhcpd: sending select for 10.0.2.15
udhcpd: lease of 10.0.2.15 obtained, lease time 86400
deleting routers
adding dns 10.0.2.3
OK

Welcome to Buildroot
buildroot login: 2020
Password:
Login incorrect
buildroot login:
buildroot login: root
#

```

## 退出qemu

ctrl+a x

```

#
# QEMU: Terminated
root@6a50e217856f:/home/oslab/lab0#

```

## 3.5 使用gdb调试内核

### Terminal 1

```

# qemu-system-riscv64 -nographic -machine virt -kernel
build/linux/arch/riscv/boot/Image \
-device virtio-blk-device,drive=hd0 -append "root=/dev/vda ro console=ttyS0" \
-bios default -drive file=rootfs.ext4,format=raw,id=hd0 \
-netdev user,id=net0 -device virtio-net-device,netdev=net0 -S -s

```

```

root@6a50e217856f:/home/oslab/lab0# qemu-system-riscv64 -nographic -machine virt -kernel build/linux/arch/riscv/boot/Image \
> -device virtio-blk-device,drive=hd0 -append "root=/dev/vda ro console=ttyS0" \
> -bios default -drive file=rootfs.ext4,format=raw,id=hd0 \
> -netdev user,id=net0 -device virtio-net-device,netdev=net0 -S -s

```

## Terminal 2

```
# riscv64-unknown-linux-gnu-gdb build/linux/vmlinux
```

```
skyyyy@maki:~$ docker exec -it 6a50 /bin/bash
root@6a50e217856f:/# cd home/oslab/lab0
root@6a50e217856f:/home/oslab/lab0# riscv64-unknown-linux-gnu-gdb build/linux/vmlinux
GNU gdb (GDB) 9.1
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "--host=x86_64-pc-linux-gnu --target=riscv64-unknown-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from build/linux/vmlinux...
(gdb) █
```

## 执行gdb命令

```
(gdb) target remote localhost:1234
```

- 含义：target remote 命令表示远程调试，而1234是默认的用于调试连接的端口号。
- 执行结果：

```
(gdb) target remote localhost:1234
Remote debugging using localhost:1234
0x0000000000000100 in ?? ()
(gdb)
```

```
(gdb) b start_kernel
(gdb) b *0x80000000
(gdb) b *0x80200000
(gdb) info breakpoints
(gdb) delete 2
(gdb) info breakpoints
```

- 含义：

`b start_kernel` 指设置断点使程序在调用start\_kernel函数时断住

`b *0x80000000` 指在地址0x80000000处设置断点

`b *0x80200000` 指在地址0x80200000处设置断点



`info breakpoints` 指打印所有已设置断点和相关信息

`delete 2` 指删除标号为2的断点

- 执行结果：

```
(gdb) b start_kernel
Breakpoint 1 at 0xffffffe000001714: file /home/oslab/lab0/linux/init/main.c, line 837.
(gdb) b *0x80000000
Breakpoint 2 at 0x80000000
(gdb) b *0x80200000
Breakpoint 3 at 0x80200000
(gdb) info breakpoints
Num      Type             Disp Enb Address                  What
1        breakpoint       keep y   0xffffffe000001714 in start_kernel
                                     at /home/oslab/lab0/linux/init/main.c:837
2        breakpoint       keep y   0x0000000080000000
3        breakpoint       keep y   0x0000000080200000
(gdb) delete 2
(gdb) delete 2
No breakpoint number 2.
(gdb) info breakpoints
Num      Type             Disp Enb Address                  What
1        breakpoint       keep y   0xffffffe000001714 in start_kernel
                                     at /home/oslab/lab0/linux/init/main.c:837
3        breakpoint       keep y   0x0000000080200000
(gdb)
```

```
(gdb) continue
(gdb) delete 3
(gdb) continue
(gdb) step
(gdb) s
(gdb) (不做输入, 直接回车)
(gdb) next
(gdb) n
(gdb) (不做输入, 直接回车)
```

- 含义：

`continue` 指从断点处继续运行程序

`delete 3` 指删除标号为3的断点

`step/s` 指单步进入函数内部并执行完

`next/n` 指单步执行下一条语句

- 执行结果：

```

(gdb) continue
Continuing.

Breakpoint 3, 0x0000000080200000 in ?? ()
(gdb) delete 3
(gdb) continue
Continuing.

Breakpoint 1, start_kernel () at /home/oslab/lab0/linux/init/main.c:837
837      set_task_stack_end_magic(&init_task);
(gdb) step
set_task_stack_end_magic (tsk=<optimized out>)
    at /home/oslab/lab0/linux/kernel/fork.c:863
863      *stackend = STACK_END_MAGIC;    /* for overflow detection */
(gdb) s
start_kernel () at /home/oslab/lab0/linux/init/main.c:838
838      smp_setup_processor_id();
(gdb)
smp_setup_processor_id () at /home/oslab/lab0/linux/arch/riscv/kernel/smp.c:38
38      cpuid_to_hartid_map(0) = boot_cpu_hartid;
(gdb) next
start_kernel () at /home/oslab/lab0/linux/init/main.c:841
841      cgroup_init_early();
(gdb) n
843      local_irq_disable();
(gdb)
844      early_boot_irqs_disabled = true;
(gdb)

```

```

(gdb) disassemble
(gdb) nexti
(gdb) n
(gdb) stepi
(gdb) s

```

- 含义：

`disassemble` 指反汇编pc附近的函数，即start\_kernel

`nexti` 指单步执行一条机器指令

`stepi` 指单步进入一条机器指令

- 执行结果：

```
(gdb) disassemble
Dump of assembler code for function start_kernel:
0xffffffff000001714 <+0>:      addi      sp,sp,-80
0xffffffff000001716 <+2>:      sd        ra,72(sp)
0xffffffff000001718 <+4>:      sd        s0,64(sp)
0xffffffff00000171a <+6>:      sd        s1,56(sp)
0xffffffff00000171c <+8>:      addi      s0,sp,80
0xffffffff00000171e <+10>:     sd        s2,48(sp)
0xffffffff000001720 <+12>:     sd        s3,40(sp)
0xffffffff000001722 <+14>:     sd        s4,32(sp)
0xffffffff000001724 <+16>:     sd        s5,24(sp)
0xffffffff000001726 <+18>:     sd        s6,16(sp)
0xffffffff000001728 <+20>:     auipc     a0,0x100a
0xffffffff00000172c <+24>:     addi      a0,a0,1560 # 0xffffffff00100bd40 <init_task>
0xffffffff000001730 <+28>:     auipc     ra,0x205
0xffffffff000001734 <+32>:     jalr      92(ra) # 0xffffffff00020678c <set_task_stack_end_magic>
0xffffffff000001738 <+36>:     jal      ra,0xffffffff000003730 <smp_setup_processor_id>
0xffffffff00000173c <+40>:     jal      ra,0xffffffff000008d4e <cgroup_init_early>
0xffffffff000001740 <+44>:     csrctl    sstatus,2
=> 0xffffffff000001744 <+48>:     li        a5,1
0xffffffff000001746 <+50>:     auipc     a4,0x106f
0xffffffff00000174a <+54>:     sb        a5,-1786(a4) # 0xffffffff00107004c <early_boot_irqs_disabled>
0xffffffff00000174e <+58>:     jal      ra,0xffffffff000004606 <boot_cpu_init>
0xffffffff000001752 <+62>:     auipc     a1,0x9ff
0xffffffff000001756 <+66>:     addi      a1,a1,-1682 # 0xffffffff000a000c0 <linux_banner>
0xffffffff00000175a <+70>:     auipc     a0,0xb3d
0xffffffff00000175e <+74>:     addi      a0,a0,-850 # 0xffffffff000b3e408
0xffffffff000001762 <+78>:     auipc     ra,0x245
0xffffffff000001766 <+82>:     jalr      -510(ra) # 0xffffffff000246564 <printk>
0xffffffff00000176a <+86>:     addi      a0,s0,-72
0xffffffff00000176e <+90>:     auipc     s4,0x106f
0xffffffff000001772 <+94>:     addi      s4,s4,-1798 # 0xffffffff001070068 <initrd_end>
0xffffffff000001776 <+98>:     jal      ra,0xffffffff0000033b4 <setup_arch>
0xffffffff00000177a <+102>:    ld        s3,0(s4)
0xffffffff00000177e <+106>:    auipc     s2,0x106f
0xffffffff000001782 <+110>:    addi      s2,s2,-1806 # 0xffffffff001070070 <initrd_start>
0xffffffff000001786 <+114>:    beqz      s3,0xffffffff0000017d4 <start_kernel+192>
0xffffffff00000178a <+118>:    addi      s1,s3,-12
0xffffffff00000178e <+122>:    li        a2,12
0xffffffff000001790 <+124>:    auipc     a1,0xb3d
0xffffffff000001794 <+128>:    addi      a1,a1,-896 # 0xffffffff000b3e410
0xffffffff000001798 <+132>:    mv        a0,s1
0xffffffff00000179a <+134>:    auipc     ra,0x47f
0xffffffff00000179e <+138>:    jalr      2030(ra) # 0xffffffff000480f88 <memcmp>
--Type <RET> for more, q to quit, c to continue without paging--
```

```
Quit
(gdb) nexti
0xffffffff000001746      844      early_boot_irqs_disabled = true;
(gdb) n
850      boot_cpu_init();
(gdb) stepi
boot_cpu_init () at /home/oslab/lab0/linux/arch/riscv/include/asm/current.h:31
31      return riscv_current_is_tp;
(gdb) s
2492      set_cpu_online(cpu, true);
(gdb)
```

- `nexti/stepi` 和 `next/step` 区别在于前者作用于机器指令层面，后者作用于源码层面；
- `next` 和 `step` 区别在于前者需要单步执行下一条指令，后者单步进入下一条指令

(gdb) continue

#由于此时无断点，continue将一直执行下去，为了退出gdb，可以：①在gdb中按住ctrl+c退出当前正在运行的gdb命令，然后再用quit退出②在qemu中先按ctrl+a，松开后再按x先退出qemu，然后在gdb中quit退出。

(gdb) quit

- 含义：

quit 退出gdb

- 执行结果：

```
2492         set_cpu_online(cpu, true);
(gdb) continue
Continuing.
^C
Program received signal SIGINT, Interrupt.
arch_cpu_idle () at /home/oslab/lab0/linux/arch/riscv/kernel/process.c:33
33         local_irq_enable();
(gdb) quit
A debugging session is active.

        Inferior 1 [process 1] will be detached.

Quit anyway? (y or n) y
Detaching from program: /home/oslab/lab0/build/linux/vmlinux, process 1
Ending remote debugging.
[Inferior 1 (process 1) detached]
root@6a50e217856f:/home/oslab/lab0#
```

## 4. 讨论和心得

---

注意容器挂载